

MIEEC

Computer Networks


Lecture note 8

Transport Layer




Transport layer services

Point-to-point: Data link vs. Transport



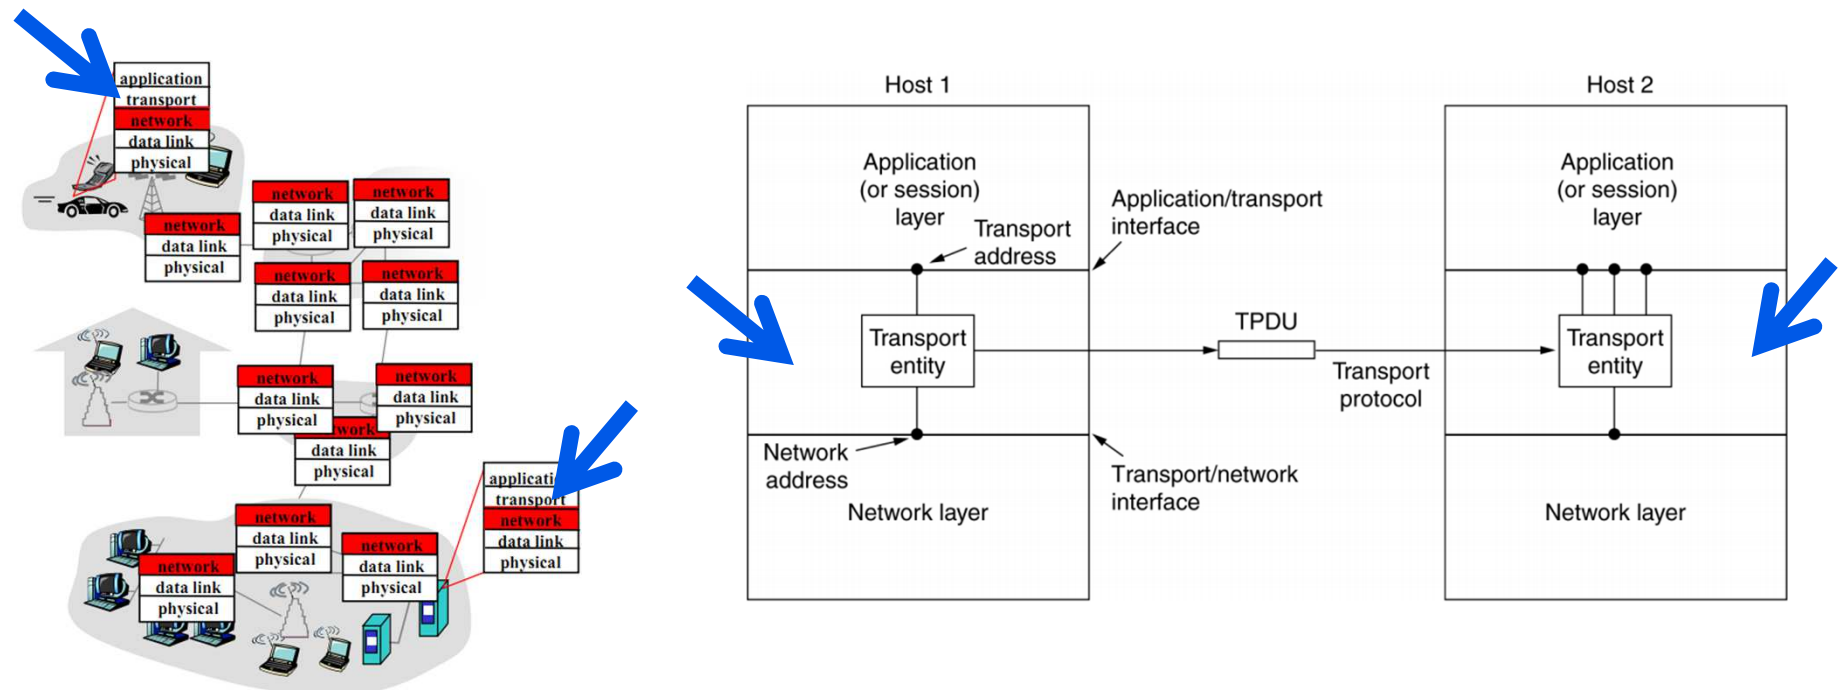
(a)



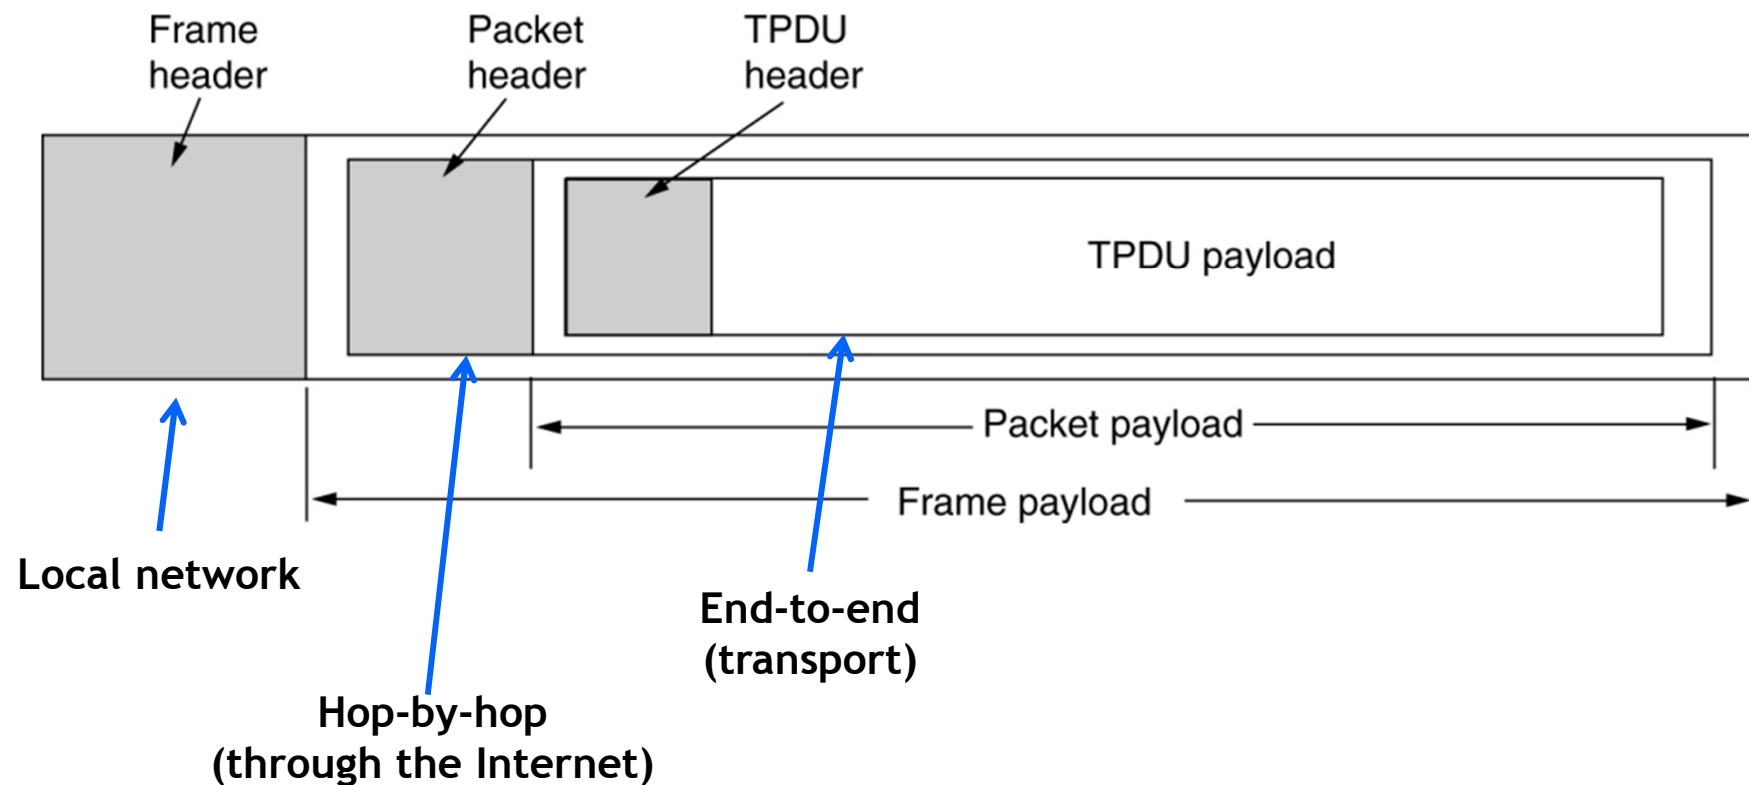
(b)

- a) P-to-P: Data link layer
- b) P-to-P: Transport layer

End-to-end service



Transport datagram encapsulating

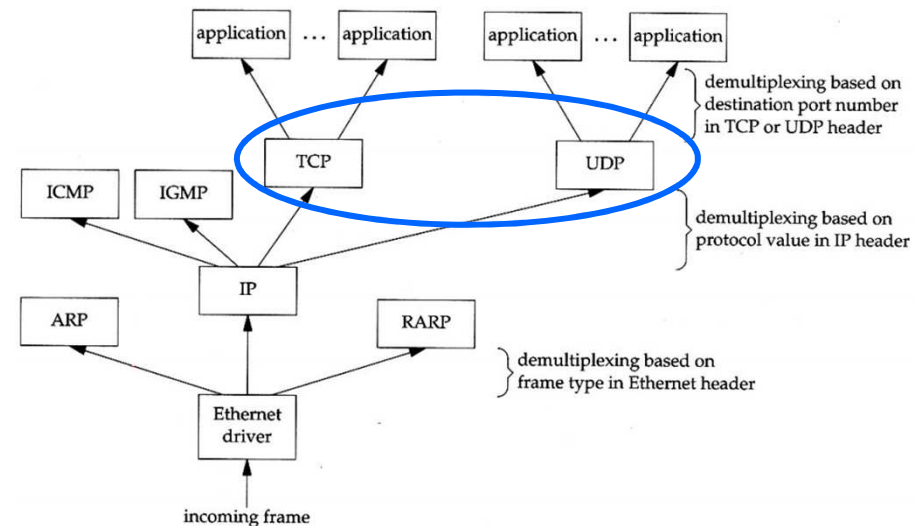


TO THINK

- What do you want out of a transport protocol?
 - Why can't you just use IP?

Demultiplexing and flows


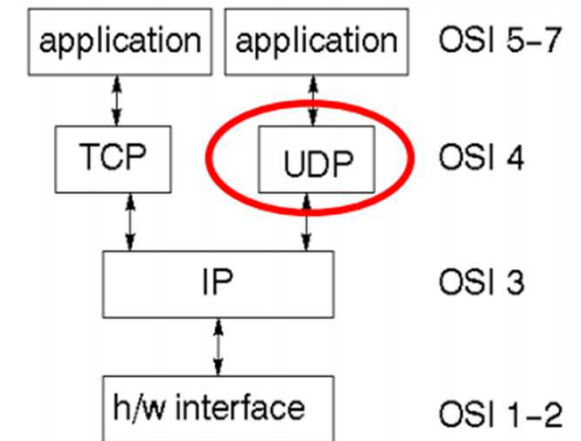
- Flow
 - Source IP
 - Destination IP
 - Source Port
 - Destination Port
 - Transport (TCP/UDP)
- 1 flow => 1 application



UDP

User Datagram Protocol

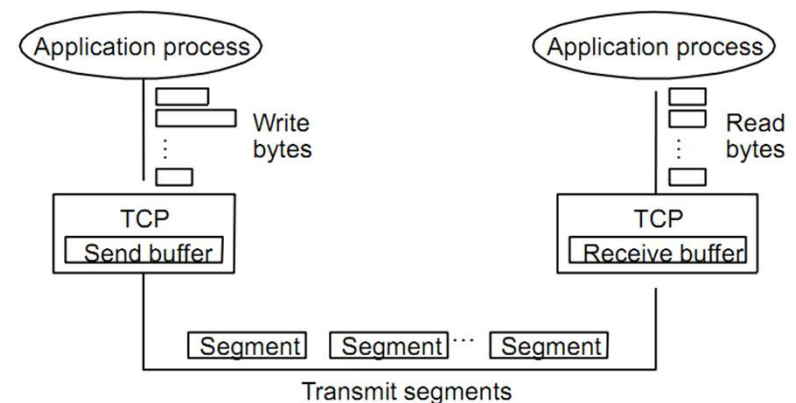
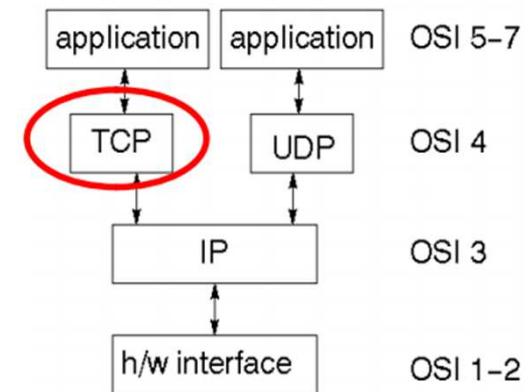
- Datagram oriented
 - Unreliable, no error control
 - Connectionless, no reordering etc.
- Provides applications with “direct”
 - Small protocol overhead
 - De-multiplexing
- UDP header
 - Port numbers identify sending/receiving process
 - UDP length : length of packet in bytes
 - Checksum covers header and data, optional



TCP

Transmission control protocol

- Connection oriented
- Full duplex
- Byte stream
- Flow control
 - Reliability
 - Avoids sending more than receiver can handle
 - ARQ
- Congestion control
 - Avoids congestion in the network

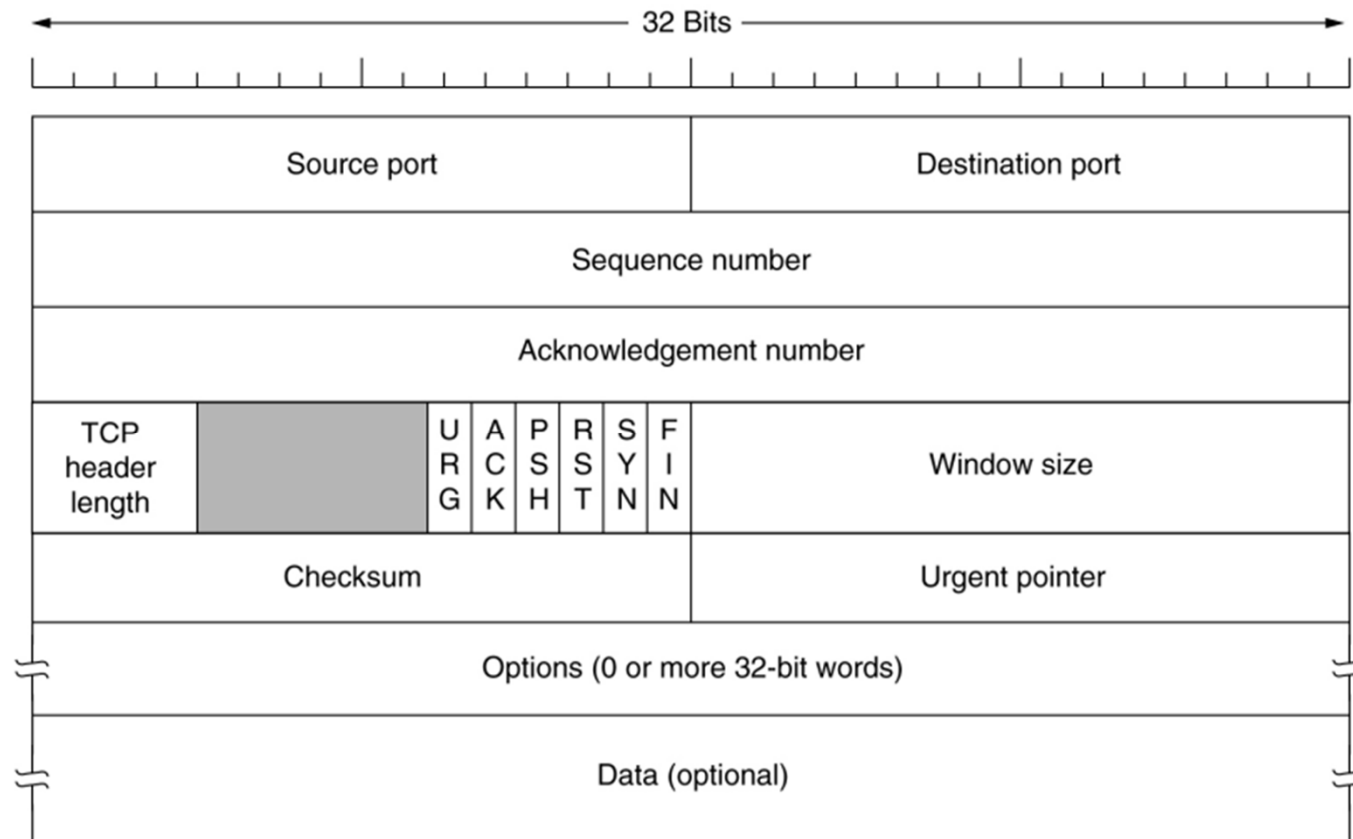


TCP

Basic TCP operation

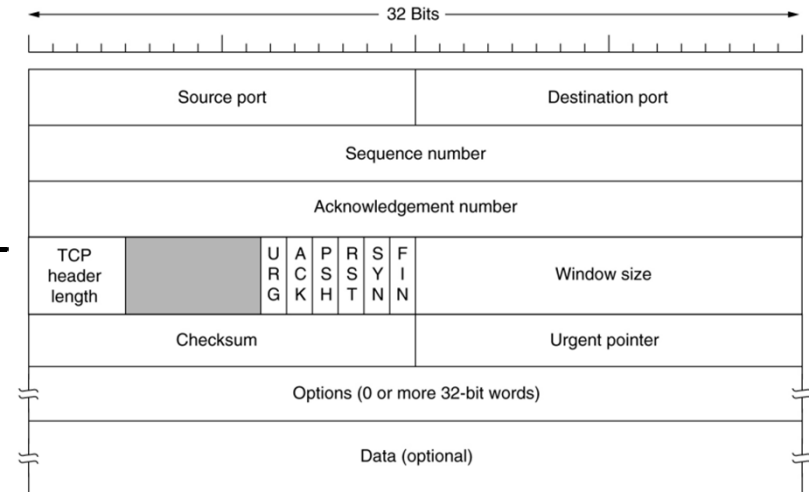
- Sender
 - Application data is broken in segments
 - TCP uses timer while waiting for an ACK for every segment
 - Un-ACKnowledged packets are retransmitted
- Receiver
 - Errors detected using checksum
 - Correctly received data is acknowledged
 - Segments are reassembled in order
 - Duplicate segments are discarded
- Window-based flow control
 - ?

TCP segment header




TCP header

- Port numbers
 - Similar to UDP, demux
- Sequence number
 - Uniquely identifies which application data is contained in the segment
 - SN in bytes, identifies 1st byte
- ACK number, piggybacking ACK's
 - Next byte the receiver is expecting
 - Implicit ACK for all bytes up to that point
- Window size
 - Flow control (ARQ) and congestion control
 - Cannot send more than window size to network
 - In bytes; can increase/decrease depending on network/traffic
- Checksum covers header and data

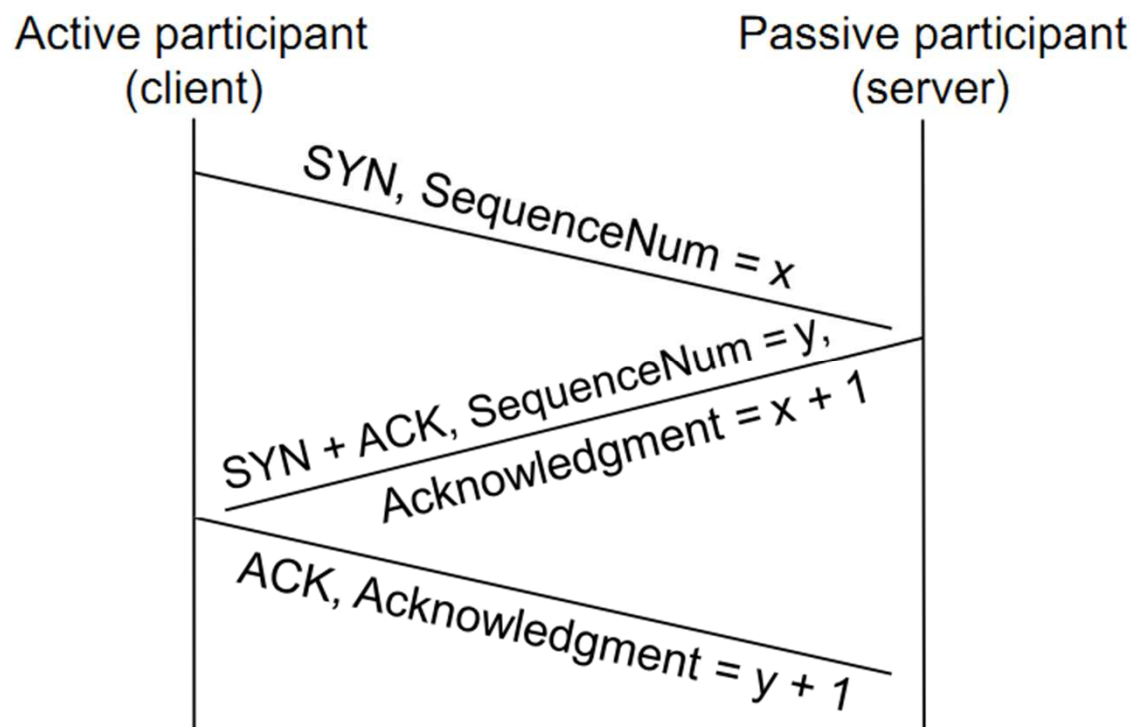


TCP sequence number

- TCP views data as streams of bytes
 - Bytes are numbered sequentially
- TCP breaks stream in segments
 - Maximum segment size MSS
- Each packet has a sequence number
 - This is the number of the 1st byte in the packet
- TCP connection is duplex
 - Different byte streams and sequence numbers in each direction

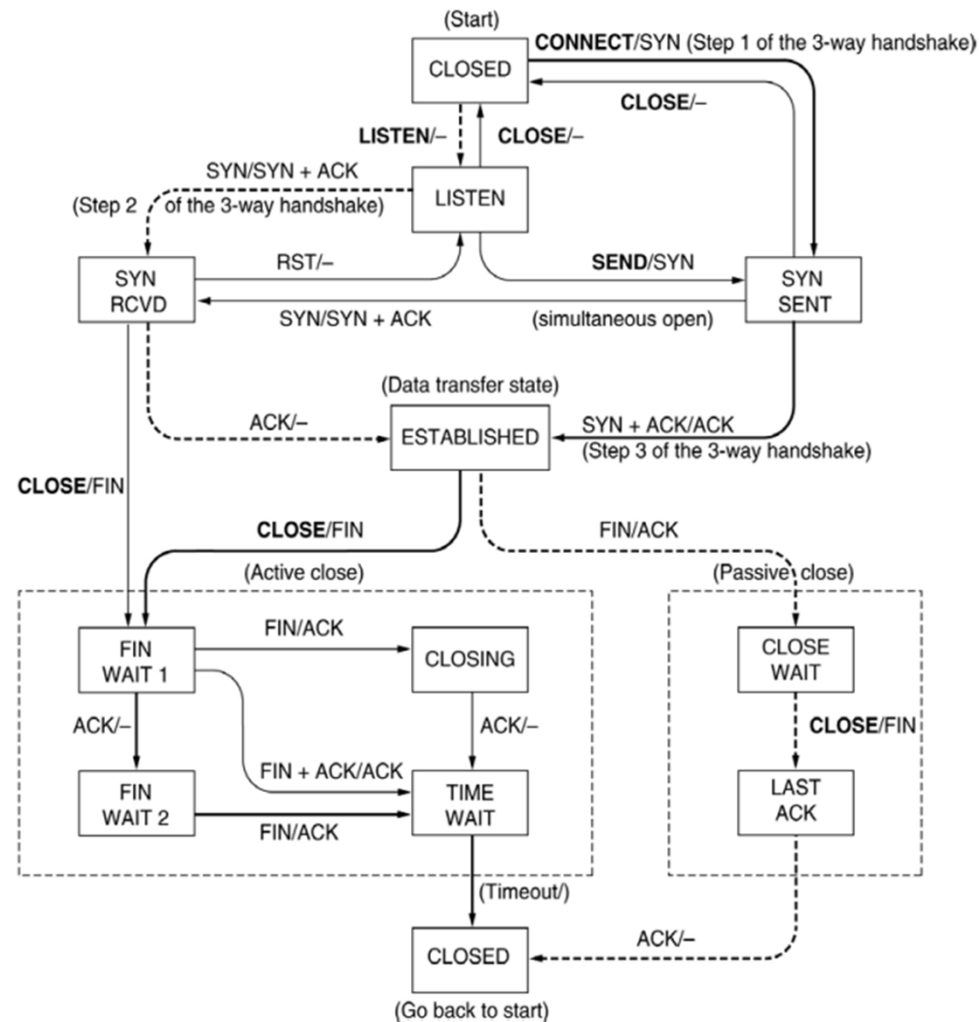


TCP connection establishment



No. .	Time	Source	Destination	Protocol	Info
13	1.246280	10.0.0.125	10.0.0.121	TCP	61432 > as-debug [SYN] Seq=0 win=8192 Len=0
14	1.246601	10.0.0.121	10.0.0.125	TCP	as-debug > 61432 [SYN, ACK] Seq=0 Ack=1 win=8192 Len=0 MSS=1460
15	1.256106	10.0.0.125	10.0.0.121	TCP	61432 > as-debug [ACK] Seq=1 Ack=1 win=8192 Len=0
16	1.263175	10.0.0.125	10.0.0.121	TCP	61432 > as-debug [PSH, ACK] Seq=1 Ack=1 win=8192 Len=16
17	1.456773	10.0.0.121	10.0.0.125	TCP	as-debug > 61432 [ACK] Seq=1 Ack=17 win=16616 Len=0
20	3.174325	10.0.0.125	10.0.0.121	UDP	Source port: domain Destination port: as-debug
24	3.314327	10.0.0.125	10.0.0.121	UDP	Source port: domain Destination port: as-debug

TCP Connection Management




Flow control

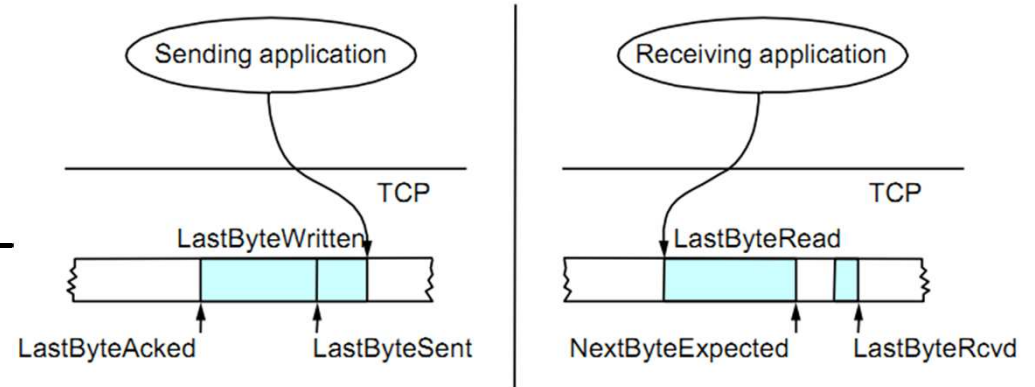
Retransmissions in TCP

A variation of Go-Back-N

- Sliding window
 - ACK contains single sequence number
 - Acknowledges all bytes with lower sequence number
- Sender retransmits single packet at a time
 - Assumes only one packet is lost (optimist)
- Error control based on byte sequences
 - not packets

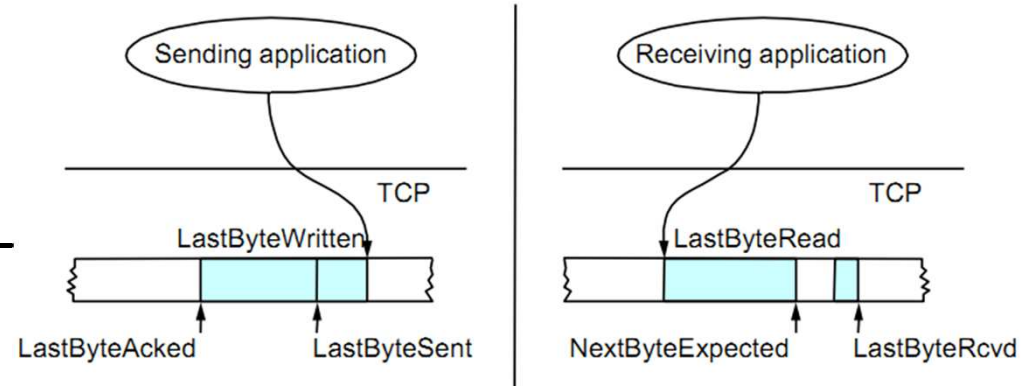


Sliding window



- **Sender**
 - $\text{LastByteAacked} \leq \text{LastByteSent}$
 - $\text{LastByteSent} \leq \text{LastByteWritten}$
 - Buffers $(\text{LastByteWritten} - \text{LastByteAacked})$ bytes
- **Receiver**
 - $\text{LastByteRead} < \text{NextByteExpected}$
 - $\text{NextByteExpected} \leq \text{LastByteRcvd} + 1$
 - Buffers $(\text{LastByteRcvd} - \text{LastByteRead})$ bytes

Flow Control



- Buffer size
 - Sender : MaxSendBuffer
 - Receiver: MaxRcvBuffer
- Receiver
 - $\text{LastByteRcvd} - \text{LastByteRead} \leq \text{MaxRcvBuffer}$
 - **AdvertisedWindow** = $\text{MaxRcvBuffer} - (\text{LastByteRcvd} - \text{LastByteRead})$
- Sender
 - $\text{LastByteWritten} - \text{LastByteAacked} \leq \text{MaxSendBuffer}$
 - $\text{LastByteSent} - \text{LastByteAacked} \leq \text{AdvertisedWindow}$
 - $\text{EffectiveWindow} = \text{AdvertisedWindow} - (\text{LastByteSent} - \text{LastByteAacked})$
- Sending application stops
 - If it needs to write y bytes and
 - $\text{LastByteWritten} - \text{LastByteAacked} + y > \text{MaxSendBuffer}$
- ACK sent when segment is received


TO THINK

- Retransmission is based on a timer.
 - How does TCP pick a timeout value?

Adaptive retransmission

- RTT, round trip time
- Measures `sampleRTT`
 - for each segment/ACK pair
- Average RTT
 - $RTT = a * RTT + (1-a) * sampleRTT$
 - a in $[0.8, 0.9]$
- Timeout: $2 * RTT$

Karn/Partridge algorithm



- sampleRTT not measured in retransmission
- Timeout doubled for each retransmission

Selective ACK

- Normal ACK
 - confirm all bytes up to this point
- Selective ACK
 - Acknowledges packets that arrive out-of-order
 - Adds bitmask of received packets
 - Implemented as a TCP option
- When to retransmit?
 - Packets may experience different delays
 - Still need to deal with reordering
 - Wait for 3 out of order packets

Congestion Control

Congestion control in TCP

- Congestion
 - More traffic than what the network can take
- Main idea:
 - Each source increases/decreases the traffic it generates
 - Based on criteria allowing
 - Flow fairness
 - Efficiency
- Approach:
 - Received ACKs regulate packet transmission

Congestion control in TCP


- Changes in channel capacity
=> adjustment of transmission rate
- New variable per connection: CongestionWindow
 - Limits the amount of traffic in the network
 - $\text{MaxWin} = \text{MIN}(\text{CongestionWindow}, \text{AdvertisedWindow})$
 - $\text{EffectiveWindow} = \text{MaxWin} - (\text{LastByteSent} - \text{LastByteAcked})$
- Goal:
 - If network congestion decreases
 - Increase CongestionWindow
 - If network congestion increases
 - Decrease CongestionWindow
- Bitrate (byte/s) => $\text{CongestionWindow} / \text{RTT}$

Congestion control in TCP

- Need to measure congestion
 - to update value of congestion window
- How does the source know about congestion?
 - By timeout
 - Wired link => low BER => low FER
 - Timeout => loss of packet
 - Packet loss => buffer in router full => drops packets => congestion
- Wireless link?


Additive Increase / Multiplicative Decrease

- Source knows about congestion by timeout
- How does it update CongestionWindow?
- Algorithm
 - Increases CongestionWindow by 1 segment
 - For each RTT => additive increase
 - Divide CongestionWindow by 2
 - When there is a packet loss => multiplicative decrease
- In practice, per received ACK
 - $\text{Increment} = \text{MSS} * (\text{MSS} / \text{CongestionWindow})$
 - $\text{CongestionWindow} += \text{Increment}$
 - MSS: Maximum segment size




Additive Increase / Multiplicative Decrease

Sawtooth wave behavior




Slow start

- At the beginning of the connection:
 - Increase CongestionWindow exponentially
 - Start with CongestionWindow=1
 - Double CongestionWindow for each ACK
- Goal
 - More quickly determine available capacity
- After first timeout
 - go to linear increase (congestion avoidance phase)



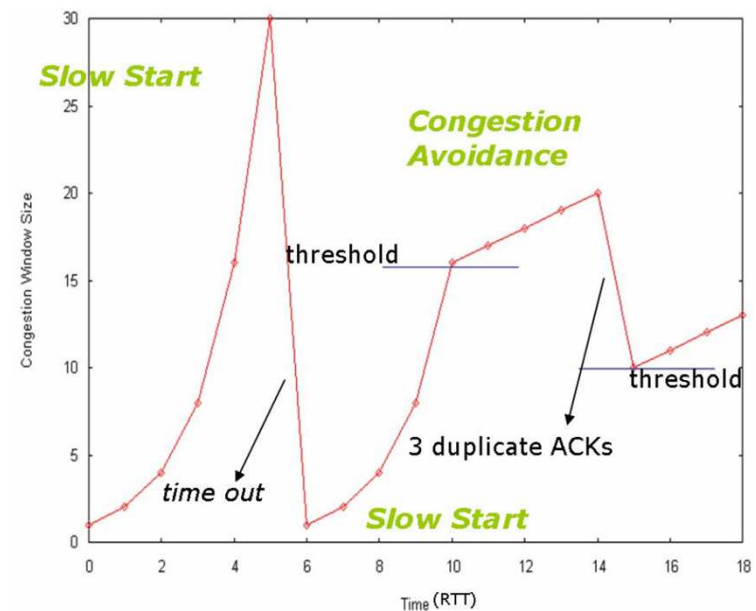

Fast retransmission, fast recovery

- If TCP timeout too large
 - Long inactivity period
 - Retransmissions delayed
- Solution
 - Fast retransmission
 - After 3 repeated ACKs




TCP - Slow start

- Slow start
 - Sender starts with $\text{CongestionWindow} = 1 \text{ sgm}$
 - Doubles CongestionWindow every RTT
- When segment loss detected **by timeout**:
 - $\text{threshold} = 1/2 \text{ CongestionWindow}$
 - $\text{CongestionWindow} = 1 \text{ sgm}$
(so that router has time to process all packets in queue)
 - Lost packet is retransmitted
 - Slow start while
 - $\text{CongestionWindow} < \text{threshold}$
 - Then Congestion Avoidance phase
 - Linear increase



TCP - Congestion Avoidance

- Additive increase
 - Increments CongestionWindow by 1 sgm per RTT
- Detection of segment loss by 3 repeated ACKs
 - Assumes packet is lost
 - Cause is not severe congestion
 - Some of the following packets have arrived
 - Retransmits lost packet
 - $\text{CongestionWindow} = \text{CongestionWindow} / 2$
 - Congestion Avoidance phase



TCP - Congestion Control

- In reality Congestion Control is a bit more complex
- RFC2581
 - “TCP Slow Start, Congestion Avoidance, Fast Retransmit, and Fast Recovery Algorithms”

-
- What services are provided by the transport layer?
 - What are the transport protocols of the TCP/IP stack?
 - What's the difference between TCP and UDP?
 - How is the connection established in TCP?
 - What's the difference between:
 - Flow control
 - Congestion control
 - How does TCP implement flow control?
 - What are the congestion control mechanisms in TCP?
 - Why is TCP congestion control so important to the Internet?

HOMEWORK

- Review slides
- Read “The Transport Layer” from:
 - Tanenbaum - Chap. 5
- Do your Moodle homework