

**MIEEC**

# **Computer Networks**

**Lecture note 9**



**Routing and shortest paths**




---

# *Graphs and shortest paths*

# Graph notation



a) Directed graph




b) Undirected graph

- $G=\{V,E\}$
- $V=\{v_1,v_2,v_3,v_4,v_5,v_6\}$ 
  - $|V| = 6$
- $E=\{(v_1,v_2), (v_2,v_3), (v_2,v_4), (v_3,v_4), (v_4,v_3), (v_3,v_5), (v_4,v_5), (v_5,v_6)\}$ 
  - $|E|=8$

- $G=\{V,E\}$
- $V=\{v_1,v_2,v_3,v_4,v_5,v_6\}$ 
  - $|V|=6$
- $E=\{(v_1,v_2), (v_2,v_3), (v_2,v_4), (v_4,v_3), (v_3,v_5), (v_4,v_5), (v_5,v_6)\}$ 
  - $|E|=7$


# Tree

---




- Tree  $T = \{V, E\}$  is a graph
  - No cycles
  - $|E| = |V| - 1$
  - Connected
    - a path exists between any pair of vertices
- Spanning tree
  - $G = \{V, E\}$
  - $T = \{V, E'\}$  is a tree and
  - $E' \subseteq E$  (is a subset of or equal to)

# Shortest path tree



- Weight on graph edges

- $G=\{V,E,w\}$
- $T=\{V,E',w\}$
- $w: E \rightarrow \mathbb{R}$




- Total cost of Tree T


- $C_{total}(T) = \sum_{i=1}^{|E'|} w(e'_i)$

- Minimum spanning tree  $T^*$


- $C_{total}(T^*) = \min_T (C_{total}(T))$
- Single MST for each graph G
- Prim, Kruskal



# Shortest path tree



- Shortest Path Tree SPT at vertex  $s$ 
  - Tree composed of
  - Union of shortest paths between  $s$  and the other vertices
  - Yes, it's a tree
  - Dijkstra, Bellman-Ford
- One SPT per vertex
  - Minimum spanning tree is not necessarily a SPT

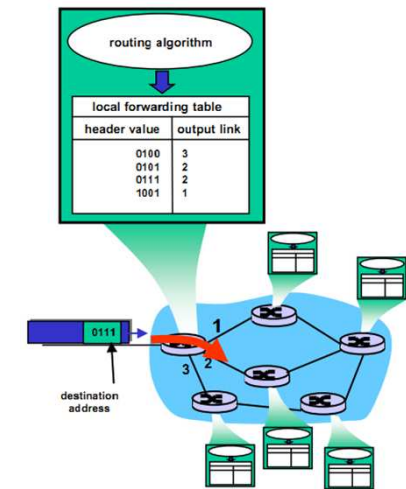


---

# *Routing in Layer 3 networks*

# Forwarding vs. Routing

- Forwarding
  - Take packets from input ports
  - Send them to output ports
  - Which output port? Forwarding table
  - DATA PLANE
- Routing
  - Compute forwarding table entries
  - By computing the (best) path packets should take
  - Distributed
    - Routers compute best paths locally
    - Send messages to each other to update routing/link information
  - CONTROL PLANE





# Importance of routing


---

- I.e. “Why do you want the best path?”
- End-to-end performance
  - Path affects QoS - delay, throughput, fairness
- Balanced used of network resources
  - Avoid congestion by choosing paths with less-loaded links
- Responsive routing
  - Compute paths on the fly
  - Responsive to link/router failures/maintenance
  - Smaller packet loss and delay

# Shortest Path Routing

---


- Path selection model:
  - Destination-based
  - Load-insensitive (link weights are static)
  - Minimum hop count / sum of link weights



# Shortest Path Problem

---

- Given network topology with link costs
  - $c(x,y)$
  - No link  $\rightarrow$  infinite cost
- Compute least-cost paths from source  $u$  to all other nodes
  - $p(v)$  : predecessor to  $v$  in the path from  $u$



# Dijkstra's Shortest Path Algorithm

---

- Iterative Algorithm
  - After  $k$  iterations know least-cost paths to  $k$  nodes
- $S$  : set of nodes whose least-cost path is known
  - Initially  $S=\{u\}$ ,  $u$  is the source node
  - Add one node to  $S$  per iteration
- $D(v)$  : current cost of path from source to node  $v$ 
  - Initially
    - $D(v) = c(u, v)$  for each node  $v$  adjacent to  $u$
    - $D(v) = \infty$  , other nodes  $v$
  - Update  $D(v)$  when shortest paths are learned





# Dijkstra's Shortest Path Algorithm

```
1 function Dijkstra(Graph, source):
2   for each vertex v in Graph:           // Initializations
3     dist[v] := infinity ;               // Unknown distance function from source to v
4     previous[v] := undefined ;          // Previous node in optimal path from source
5   end for ;
6   dist[source] := 0 ;                   // Distance from source to source
7   Q := the set of all nodes in Graph ; // All nodes in the graph are unoptimized - thus are in Q
8   while Q is not empty:                 // The main loop
9     u := vertex in Q with smallest distance in dist[] ;
10    if dist[u] = infinity:
11      break ;                           // all remaining vertices are inaccessible from source
12    end if ;
13    remove u from Q ;
14    for each neighbor v of u:           // where v has not yet been removed from Q.
15      alt := dist[u] + dist_between(u, v) ;
16      if alt < dist[v]:                 // Relax (u,v,a)
17        dist[v] := alt ;
18        previous[v] := u ;
19        decrease-key v in Q;           // Reorder v in the Queue
20      end if ;
21    end for ;
22  end while ;
23  return dist[] ;
24 end Dijkstra.
```





## Retrieve path from source to target

```
1 S := empty sequence
2 u := target
3 while previous[u] is defined:
4   insert u at the beginning of S
5   u := previous[u]
6 end while ;
```

# Dijkstra's example (1)




# Dijkstra's example (2)



# Forwarding table at u

- Shortest Path Tree



	link
v	(u,v)
w	(u,w)
x	(u,w)
y	(u,v)
z	(u,v)
s	(u,w)
t	(u,w)



# TO THINK

---

- How does router  $u$  know the link costs of non-adjacent routers?

# Link state routing


---

- Each router keeps track of its incident links
  - Link up/down
  - Link cost
- Each router broadcasts the link state
  - Every router gets a full view of the network graph
- Each router runs Dijkstra's to
  - Compute shortest paths
  - Construct forwarding table
- Example protocols
  - OSPF, Open Shortest Path First
  - IS-IS, Intermediate System - Intermediate System

# Detection of topology changes

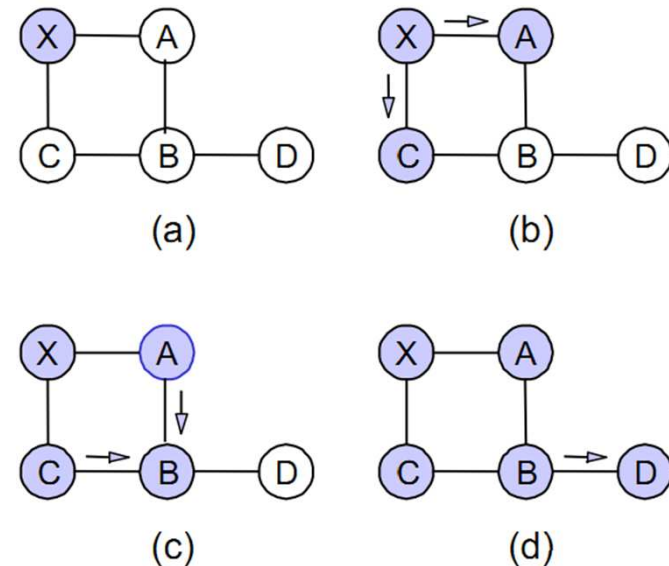
---

- Beacons generated by routers on their links
  - Periodic “hello” messages on both directions
  - A few missed “hellos” => link down



# Broadcasting link state


- How to flood the link state?
  - Every node sends link-state information through their links
  - Next nodes forward link-state info on all links
    - except the one on which the info arrived
- When to start flooding?
  - Upon topology change
    - Link failure/recovery
    - Link cost change
  - Periodically
    - Refresh link-state information
    - typically 30 min



# Scalability of link-state routing

---

- Overhead of link-state routing
  - Flooding link-state info on all the network
  - Running Dijkstra's with all nodes
- OSPF “areas” introduce hierarchy
  - Split the network, the flooding, and routing computation



# TO THINK


---

- Link-state/Dijkstra's requires knowledge of network topology at every router
  - $c(x,y)$
- Is it possible to compute shortest paths without knowing the whole topology of the network?
  - Is the distance to other nodes enough?

# Bellman-Ford Algorithm

---

- Define distances at each node  $x$ 
  - $d_x(y)$  : cost of lest-cost path from  $x$  to  $y$
- Update distances based on neighbors
  - $d_x(y) = \min(c(x, v) + d_v(y))$ , on all neighbors  $v$



# Distance Vector Algorithm

---

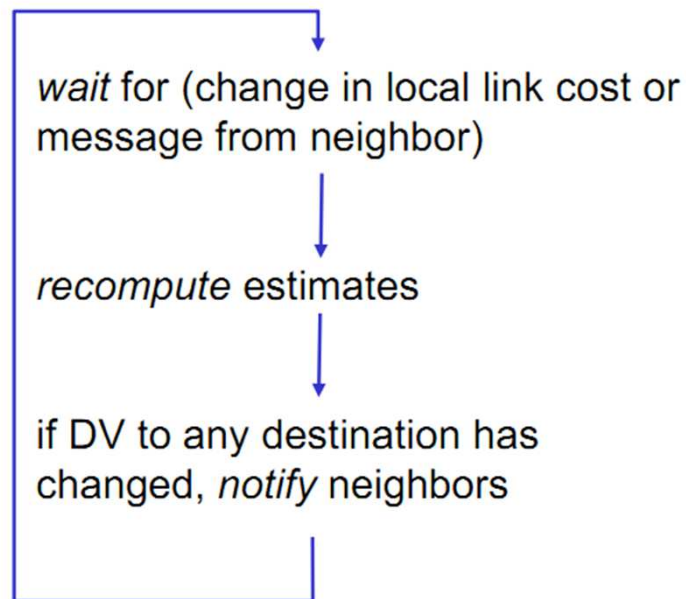
- $c(x,v)$  = cost for direct link  $(x,v)$ 
  - Node  $x$  keeps costs of direct links on  $x$
- $D_x(y)$  = estimate of least cost from  $x$  to  $y$ 
  - Node  $x$  keeps vector  $\mathbf{D}_x = [D_x(y): y \in N]$
- Node  $x$  also keeps estimate of neighbors distance vectors
  - For each neighbor  $v$ ,  $x$  keeps  $\mathbf{D}_v = [D_v(y): y \in N]$
- Each node  $v$  periodically sends  $\mathbf{D}_v$  to neighbors
  - And neighbors update their distance vectors
  - $D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\}$  for each node  $y \in N$
- Over iterations (time) vector  $\mathbf{D}_x$  converges



# Distance Vector Algorithm

---

Each node:



- Iterative
  - Each local iteration is caused by
    - Local link cost change
    - DV update message from neighbor
- Distributed
  - Node notifies neighbors when DV changes
  - Neighbors notify neighbors, if necessary

# DV Example (Initialization)

Table for A			Table for B		
Dst	Cst	Hop	Dst	Cst	Hop
A	0	A	A	4	A
B	4	B	B	0	B
C	$\infty$	–	C	$\infty$	–
D	$\infty$	–	D	3	D
E	2	E	E	$\infty$	–
F	6	F	F	1	F




Table for C			Table for D			Table for E			Table for F		
Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop
A	$\infty$	–	A	$\infty$	–	A	2	A	A	6	A
B	$\infty$	–	B	3	B	B	$\infty$	–	B	1	B
C	0	C	C	1	C	C	$\infty$	–	C	1	C
D	1	D	D	0	D	D	$\infty$	–	D	$\infty$	–
E	$\infty$	–	E	$\infty$	–	E	0	E	E	3	E
F	1	F	F	$\infty$	–	F	3	F	F	0	F

# DV Example (Step 1)

Table for A			Table for B		
Dst	Cst	Hop	Dst	Cst	Hop
A	0	A	A	4	A
B	4	B	B	0	B
C	7	F	C	2	F
D	7	B	D	3	D
E	2	E	E	4	F
F	5	E	F	1	F




Table for C			Table for D			Table for E			Table for F		
Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop
A	7	F	A	7	B	A	2	A	A	5	B
B	2	F	B	3	B	B	4	F	B	1	B
C	0	C	C	1	C	C	4	F	C	1	C
D	1	D	D	0	D	D	$\infty$	-	D	2	C
E	4	F	E	$\infty$	-	E	0	E	E	3	E
F	1	F	F	2	C	F	3	F	F	0	F

# DV Example (Step 2)

Table for A			Table for B		
Dst	Cst	Hop	Dst	Cst	Hop
A	0	A	A	4	A
B	4	B	B	0	B
C	6	E	C	2	F
D	7	B	D	3	D
E	2	E	E	4	F
F	5	E	F	1	F




Table for C			Table for D			Table for E			Table for F		
Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop
A	6	F	A	7	B	A	2	A	A	5	B
B	2	F	B	3	B	B	4	F	B	1	B
C	0	C	C	1	C	C	4	F	C	1	C
D	1	D	D	0	D	D	5	F	D	2	C
E	4	F	E	5	C	E	0	E	E	3	E
F	1	F	F	2	C	F	3	F	F	0	F

# RIP

## Routing Information Protocol


---

- Distance Vector protocol
  - Nodes send distance vectors to neighbors every 30 seconds
  - Or when an update message causes change in routing
- RIP is limited to small networks

# BGP

## Exterior gateway routing protocol

- Path vector protocol
- Internet routing



Information F receives  
from its neighbors about D

From B: "I use BCD"

From G: "I use GCD"

From I: "I use IFGCD"

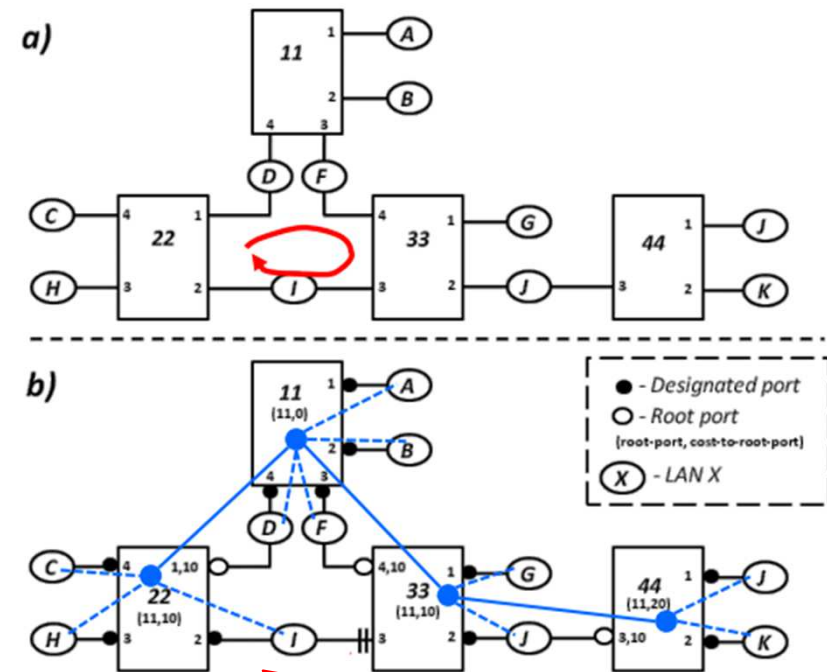
From E: "I use EFGCD"

---

# *Unique spanning trees in Ethernet networks*

# Networks of Layer 2 switches

- Ethernet frame
  - No hop-count
  - Could loop forever
  - Broadcast frame, misconfiguration
- Layer 2 network
  - Requires tree topology
  - Single path between any station pair
- Spanning Tree Protocol
  - Runs in bridges
  - Helps building the spanning tree
  - Blocks ports






# Constructing a Spanning Tree

---

- Distributed algorithm
  - Switches need to elect a “root” for the tree
    - The switch with the smallest identifier
  - Each switch finds out if its interface is on the shortest path from the root
  - Messages (Y,d,X)
    - From node X
    - Claiming node Y is the root
    - And the distance is d



# Steps in ST algorithm


---

- Initially each switch thinks it's the root
  - Sends message out on every interface
  - Identifies itself as root
  - Switch X announces (X,0,X)
- Other switches update their view of the root
  - Upon receiving a message, check root id
  - If the new id is smaller, start viewing that switch as root
- Switches compute their distance from the root
  - Add 1 to the distance received from neighbor
  - Identify interfaces not in the shortest path to root
  - Exclude these interfaces from the spanning tree

# Example - Switch #4 viewpoint

---


- Switch #4 thinks it's the root
  - Sends (4,0,4) to 2 and 7
- Then switch #4 hears from #2
  - Receives (2,0,2)
  - Thinks #2 is the root
  - Realizes it's 1 hop from the root
- Switch #4 hears from #7
  - Receives (2,1,7) from #7
  - Realizes this is a longer path
  - Prefers its 1 hop path to #2
  - Removes 4-7 link from tree



# Example - Switch #4 viewpoint

---

- Switch #2 hears about #1
  - Hears (1,1,3) from #3
  - Starts treating #1 as root
  - Sends (1,2,2) to neighbors
- Then switch #4 hears from #2
  - Starts treating #1 as root
  - Sends (1,3,4) to neighbors
- Switch #4 hears from #7
  - Receives (1,3,7) from #7
  - Realizes this is a longer path
  - Prefers its 3 hop path to #1
  - Removes 4-7 link from tree




---

# *Maximum flow in a network*

# Flow network model

---

- Flow network
  - Source s
  - Sink t
  - Nodes a,b,c
- Edges labeled with capacities
  - (e.g. bit/s)
- Communication networks are not flow networks
  - They are queue networks
  - Flow network approach used to find limit values



# Maximum capacity of a flow network

---





- Max-flow min-cut theorem
  - Maximum amount of flow transferrable through a network
  - Equals minimum value of all simple cuts in the network
- Cut
  - Split of nodes  $V$  in  $G=\{V,E\}$  into disjoint sets  $S$  and  $T$
  - $S \cup T = V$
  - There are  $2^{|V|-2}$  possible cuts
- Capacity of cut  $(S,T)$ 
  - $c(S,T) = \sum_{(u,v)|u \in S, v \in T, (u,v) \in E} c(u,v)$

# Max-flow Min-cut example

Max-flow = 10

$2^{|5|-2} = 8$  possible cuts

Cut	Vertices					$c(S, T)$	Feasibility
	s	a	b	c	t		
1	S	S	S	S	T	10	✓
2	S	S	S	T	T	19	✓
3	S	S	T	S	T	13	✓
4	S	S	T	T	T	17	✓
5	S	T	S	S	T	-	×
6	S	T	S	T	T	-	×
7	S	T	T	S	T	11	✓
8	S	T	T	T	T	12	✓





- 
- What is a graph?
  - What is a spanning tree?
  - What is a shortest path tree?
  - How are paths defined in a network?
  - How does Dijkstra's algorithm work?
  - How does a link-state algorithm work?
  - How do nodes find out about their neighbors?
  - How does the Bellman-Ford algorithm work?
  - How does a Distance-Vector algorithm work?
  - What are the limitations of a layer 2 network of switches?
  - How does the IEEE Spanning Tree protocol work?
  - What is the maximum capacity of a flow network?

# **HOMEWORK**

---

- Review slides
- Read from Tanenbaum
  - Section 5.2 - Routing Algorithms
  - Section 4.7.3 - Spanning Tree Bridges
- Do your Moodle homework